



(a) One full-attention layer,  $Q = 3$



(b) Effective KV-cache bandwidth



(c) End-to-end, by attention backend

